

# Statistics 360: Advanced R for Data Science

## Lecture 2

Brad McNeney

# R Data Structures

Vectors

Matrices, data frames and tibbles

Logical and relational operators

Aside: Special values

# Reading

- ▶ Text, chapters 3 and 4

# R Data Structures

# R Data Structures

- ▶ Fundamentally, all common data structures in R are vectors, which can be “atomic” or “list”.
- ▶ R has no true scalars; e.g., in `x<-1`, `x` is a vector of length one.
- ▶ Use `str()` to see the structure of an object

# Types of objects

- ▶ All R objects have a “type”, that describes how it is stored in computer memory.
- ▶ Common types we will encounter are “logical”, “integer”, “double”, “character” and “list”.
  - ▶ Find the type of an object with `typeof()`.

```
x <- 6 # stores as double by default  
typeof(x)
```

```
## [1] "double"
```

```
y <- 6L # The "L" suffix forces storage as integer  
typeof(y)
```

```
## [1] "integer"
```

# Type *versus* Mode

- ▶ In addition to the type of an object, there is its “mode”.
- ▶ The mode of an object is generally the same as its type, but the modes are coarser.
  - ▶ For example, integer and double types are both of mode “numeric”.
- ▶ I don’t understand the need for mode.
  - ▶ The only reason I mention it is that the `str()` function sometimes reports the mode of an object, rather than its type, so we will frequently see reports of numeric objects.

```
mode(x)
```

```
## [1] "numeric"
```

```
mode(y)
```

```
## [1] "numeric"
```

# Vectors



# Vectors

- ▶ Vectors can be either atomic or list
  - ▶ The elements of an atomic vector must be the same type.
  - ▶ Lists can be comprised of multiple data types
- ▶ Empty vectors can be created by the `vector()` function:

```
# help("vector")  
avec <- vector(mode="numeric",length=4)  
lvec <- vector(mode="list",length=4)
```

- ▶ Data vectors can be created with `c()` or `list()`:

```
avec <- c(54,210,77)  
lvec <- list(54,210,77,c("grey","thin"))
```

# Combining vectors

- Use `c()` to combine vectors

```
c(avec,c(100,101))
```

```
## [1] 54 210 77 100 101
```

```
c(lvec,TRUE)
```

```
## [[1]]
```

```
## [1] 54
```

```
##
```

```
## [[2]]
```

```
## [1] 210
```

```
##
```

```
## [[3]]
```

```
## [1] 77
```

```
##
```

```
## [[4]]
```

```
## [1] "grey" "thin"
```

```
##
```

```
## [[5]]
```

```
## [1] TRUE
```

## Vector attributes

- ▶ Vectors have a type and length and, optionally, attributes such as names.
  - ▶ As we have seen, we find the type of an object with `typeof()`.
  - ▶ Find the length of a vector with `length()`.

```
typeof(avec)
```

```
## [1] "double"
```

```
length(avec)
```

```
## [1] 3
```

```
str(avec)
```

```
##  num [1:3] 54 210 77
```

```
typeof(lvec)
```

```
## [1] "list"
```

```
length(lvec)
```

```
## [1] 4
```

```
names(lvec) = c("age", "weight", "height", "hair")  
str(lvec)
```

```
## List of 4
```

```
## $ age : num 54
```

```
## $ weight: num 210
```

```
## $ height: num 77
```

```
## $ hair : chr [1:2] "grey" "thin"
```

► We can specify element names when creating a vector; e.g.:

```
lvec <- list(age=54, weight=210, height=77, hair=c("grey", "thin"))
```

## More on attributes

- ▶ More generally, attributes are used to store meta-data on an object.
- ▶ Get with `attributes()` and set with `attr()` or `structure()`.

```
attributes(lvec)
```

```
## $names  
## [1] "age"      "weight" "height" "hair"  
  
attr(lvec,"at1") <- "abc" # add by name  
attr(lvec,"at2") <- 1:3  
attributes(lvec)
```

```
## $names  
## [1] "age"      "weight" "height" "hair"  
##  
## $at1  
## [1] "abc"  
##  
## $at2  
## [1] 1 2 3
```

```
x <- structure(list(1,2,3),names=c("a","b","c"),at1="abc",at2=1:3)
attributes(x)
```

```
## $names
## [1] "a" "b" "c"
##
## $at1
## [1] "abc"
##
## $at2
## [1] 1 2 3
```

# Attributes after operations

- ▶ Tend to be dropped. If not, they may not make sense.

```
x <- structure(1:3,at1="abc")  
y <- structure(3:1,at1="def")  
x
```

```
## [1] 1 2 3  
## attr(,"at1")  
## [1] "abc"
```

```
y
```

```
## [1] 3 2 1  
## attr(,"at1")  
## [1] "def"
```

```
x[1:2]
```

```
## [1] 1 2
```

```
x*y
```

```
## [1] 3 4 3  
## attr(,"at1")  
## [1] "abc"
```

## The class attribute

- ▶ As special attribute named `class` is a key component of the “S3” object-oriented programming system (more on this later).
- ▶ Get and set with the `class()` function.

```
class(lvec) <- "prof"  
lvec
```

```
## $age  
## [1] 54  
##  
## $weight  
## [1] 210  
##  
## $height  
## [1] 77  
##  
## $hair  
## [1] "grey" "thin"  
##  
## attr(,"at1")  
## [1] "abc"  
## attr(,"at2")  
## [1] 1 2 3  
## attr(,"class")  
## [1] "prof"
```



## S3 atomic vectors

- ▶ There are four important classes of atomic vector: factors, Dates, POSIXct and difftime.
- ▶ We will discuss factors. See the text, section 3.4 for information on the others.

# Factors

- ▶ The statistical concept of a factor is important in experimental design.
- ▶ Factors are implemented in R as atomic vectors with attributes `class` and `levels`:

```
trt <- factor(c("drug1","placebo","placebo","drug2"))  
attributes(trt)
```

```
## $levels  
## [1] "drug1"   "drug2"   "placebo"  
##  
## $class  
## [1] "factor"
```

```
str(trt)
```

```
## Factor w/ 3 levels "drug1","drug2",...: 1 3 3 2
```

- ▶ The levels are coded numerically (1, 2 and 3) with assigned labels ordered alphabetically ("drug1", "drug2" and "placebo") by default.
- ▶ You can specify an order to the factors with the `level` argument:

```
trt <- factor(c("drug1", "placebo", "placebo", "drug2"),  
              levels=c("placebo", "drug1", "drug2"))  
trt
```

```
## [1] drug1  placebo placebo drug2  
## Levels: placebo drug1 drug2
```

# Subsetting vectors and extracting elements

- ▶ Subset with `[]` or by name:

```
lvec[c(1,3)] # same as lvec[c("age", "height")]
```

```
## $age  
## [1] 54  
##  
## $height  
## [1] 77
```

- ▶ Extract individual elements with `[[`, or `$` for named objects:

```
lvec[[4]]
```

```
## [1] "grey" "thin"
```

```
lvec$hair
```

```
## [1] "grey" "thin"
```

## Take care using `[[` for lists

- ▶ `[[ ]]` (and `$`) can only return one item at a time. Trying to extract more than one has unexpected results.

```
lvec[[c(4,2)]] # the [[2]] element of lvec[[4]] !!
```

```
## [1] "thin"
```

```
# Error when you do this: lvec$c("hair","weight")
```

# Subsetting factors

- ▶ Subsetting may remove all instances of a level, but the level will be retained in the data structure

```
trt[1:3]
```

```
## [1] drug1    placebo placebo  
## Levels: placebo drug1 drug2
```

- ▶ If subsetting is intended to remove a level of the factor, use `drop=TRUE`

```
trt[1:3,drop=TRUE]
```

```
## [1] drug1    placebo placebo  
## Levels: placebo drug1
```

# Subsetting and assignment

- ▶ You can combine subsetting and assignment to change the value of vectors.

```
avec
```

```
## [1] 54 210 77
```

```
avec[2] <- 220  
avec
```

```
## [1] 54 220 77
```

# Assignment to vector elements

- ▶ To assign to a vector element, it is clearer to use `[[` rather than `[`.
- ▶ Also, for lists, assignment with `[` requires that the replacement element be of length 1; `[[` does not have this restriction

```
lvec[3:4] <- c("Hi", "there")  
lvec[3:4]
```

```
## $height  
## [1] "Hi"  
##  
## $hair  
## [1] "there"
```



```
lvec[4] <- c("All", "of", "this")
```

```
## Warning in lvec[4] <- c("All", "of", "this"): number of items to replace is  
## a multiple of replacement length
```

```
lvec[4] # Only used first element of replacement vector
```

```
## $hair  
## [1] "All"
```

```
lvec[[4]] <- c("All", "of", "this")  
lvec[3:4]
```

```
## $height  
## [1] "Hi"  
##  
## $hair  
## [1] "All" "of" "this"
```

# Coercion: atomic vectors to lists

- ▶ Atomic vectors can be coerced to lists with `as.list()`:

```
avec = c(age=54,weight=210,height=77)
avec
```

```
##      age weight height
##      54     210      77
```

```
as.list(avec)
```

```
## $age
## [1] 54
##
## $weight
## [1] 210
##
## $height
## [1] 77
```

## Coercion: lists to atomic vectors

- ▶ Lists can be “flattened” into atomic vectors with `unlist()`:

```
unlist(lvec)
```

```
##      age weight height  hair1  hair2  hair3  
##    "54"  "210"   "Hi"  "All"   "of" "this"
```

- ▶ Notice how the numeric values are coerced to the more flexible character type.
- ▶ The order of flexibility, from least to most, is logical, integer, numeric, character.

## Coercion: factors to atomic vectors

- ▶ We saw how to use `factor()` to coerce an atomic vector to a factor.
- ▶ Use `as.vector()` to coerce a factor back to an atomic vector.
- ▶ The result is a character vector. You may need to use `as.numeric()` to coerce to numeric, if required.

```
a <- factor(c(2,1,1,2))  
as.vector(a)
```

```
## [1] "2" "1" "1" "2"
```

```
as.numeric(as.vector(a))
```

```
## [1] 2 1 1 2
```

## Matrices, data frames and tibbles

# Matrices and data frames

- ▶ Matrices are implemented as atomic vectors with a “dim” attribute of length 2 (number of rows, number of columns).
- ▶ As an atomic vector, elements of a matrix must all be of the same type.
- ▶ Data frames are lists where each list element has the same length. Thus data frames can include columns of varying type.

# Matrices

- ▶ Matrices can be created with the `matrix()` function as in

```
A <- matrix(1:4,nrow=2,ncol=2)
```

```
A
```

```
##      [,1] [,2]
```

```
## [1,]    1    3
```

```
## [2,]    2    4
```

- ▶ Here `1:4` is the same as `c(1,2,3,4)`

- ▶ The default is to read the data vector into the matrix column-by-column. To read row-by-row instead use the `byrow=TRUE` argument:

```
A <- matrix(1:4,nrow=2,ncol=2,byrow=TRUE)
```

```
A
```

```
##      [,1] [,2]
```

```
## [1,]    1    2
```

```
## [2,]    3    4
```

```
A <- matrix(c( 1, 2,  
              3, 4),ncol=2,nrow=2,byrow=TRUE)
```



# Combining matrices

- Combine matrices with `rbind()` and `cbind()`:

```
rbind(A,matrix(c(5,6),nrow=1,ncol=2))
```

```
##      [,1] [,2]  
## [1,]    1    2  
## [2,]    3    4  
## [3,]    5    6
```

```
cbind(A,A)
```

```
##      [,1] [,2] [,3] [,4]  
## [1,]    1    2    1    2  
## [2,]    3    4    3    4
```

# Matrix attributes

- ▶ Matrices have a type, dimension and optional attributes such as dimnames (row and column names).

```
typeof(A)
```

```
## [1] "double"
```

```
dim(A)
```

```
## [1] 2 2
```

```
colnames(A) <- c("var1","var2")
rownames(A) <- c("subj1","subj2")
A
```

```
##      var1 var2
## subj1    1    2
## subj2    3    4
```

```
str(A)
```

```
##  num [1:2, 1:2] 1 3 2 4
## - attr(*, "dimnames")=List of 2
##  ..$ : chr [1:2] "subj1" "subj2"
##  ..$ : chr [1:2] "var1" "var2"
```

# Subsetting matrices

- ▶ Subset with [ and a comma to separate rows from columns:

```
A[1,1]
```

```
## [1] 1
```

```
A[1,]
```

```
## var1 var2
```

```
##    1    2
```

```
A[,1]
```

```
## subj1 subj2
```

```
##     1     3
```

- ▶ When a subsetting operation leads to a vector, the dimension of the object is “dropped” from 2 to 1. To prevent this use drop=FALSE:

```
A[1,,drop=FALSE]
```

```
##      var1 var2
```

```
## subj1    1    2
```

# Extracting elements from matrices

- ▶ Can use `[[` to extract elements, but this is not necessary because of the way subsetting to a single element drops to a vector of length 1 by default:

```
A[[1,1]]
```

```
## [1] 1
```

```
A[1,1]
```

```
## [1] 1
```

## Coercion: Matrices to/from vectors

- ▶ We have already seen how `matrix()` coerces a vector to a matrix
- ▶ `as.vector()` applied to a matrix removes the `dim` attribute and creates a vector by concatenating columns:

```
as.vector(A)
```

```
## [1] 1 3 2 4
```

# Data frames

- ▶ Data frames (class `data.frame`) are the usual way to store data in R.
  - ▶ Rows are intended to be observational units, columns variables
  - ▶ Implemented as a list (columns are list elements), but also behave like a matrix in terms of combining and subsetting.
- ▶ Create with `data.frame`:

```
set.seed(1)
n <- 4
x <- 1:n; y <- rnorm(n,mean=x,sd=1) # multiple commands separated by ;
dd <- data.frame(x=x,y=y) # like making a list
str(dd)
```

```
## 'data.frame':    4 obs. of  2 variables:
## $ x: int  1 2 3 4
## $ y: num  0.374 2.184 2.164 5.595
```

# Row and column names

- ▶ Get and set with `rownames()` and `colnames()`, respectively.
- ▶ `names()` also gets and sets the column names

```
colnames(dd); names(dd)
```

```
## [1] "x" "y"
```

```
## [1] "x" "y"
```

```
rownames(dd) <- paste0("subj", 1:nrow(dd))  
rownames(dd)
```

```
## [1] "subj1" "subj2" "subj3" "subj4"
```



## Subsetting and combining data frames like a list

```
dd$x
```

```
## [1] 1 2 3 4
```

```
dd[[1]]
```

```
## [1] 1 2 3 4
```

## Subsetting and combining data frames like a matrix

```
dd[1:2,]
```

```
##           x           y
## subj1 1 0.3735462
## subj2 2 2.1836433
```

```
zz = data.frame(z=runif(4))
cbind(dd,zz)
```

```
##           x           y           z
## subj1 1 0.3735462 0.62911404
## subj2 2 2.1836433 0.06178627
## subj3 3 2.1643714 0.20597457
## subj4 4 5.5952808 0.17655675
```

# tibbles

- ▶ A tidyverse replacement (or extension) of the data frame with different default behaviour.
- ▶ See section 3.6 for a discussion of some of the differences.

```
library(tibble)
tt <- tibble(x = 1:3, y = letters[1:3])
typeof(tt)
```

```
## [1] "list"
```

```
attributes(tt)
```

```
## $class
## [1] "tbl_df"      "tbl"        "data.frame"
##
## $row.names
## [1] 1 2 3
##
## $names
## [1] "x" "y"
```

## More on subsetting

- ▶ one-dimensional subsets of data frames are coerced to atomic vectors, but not so with tibbles.

```
dd[,1]; dd[1,1]
```

```
## [1] 1 2 3 4
```

```
## [1] 1
```

```
tt[,1]; tt[1,1]
```

```
## # A tibble: 3 x 1
```

```
##       x
```

```
##   <int>
```

```
## 1     1
```

```
## 2     2
```

```
## 3     3
```

```
## # A tibble: 1 x 1
```

```
##       x
```

```
##   <int>
```

```
## 1     1
```

# List columns

- ▶ One way that tibbles improve on data frames is easier handling of columns whose elements are lists.

```
tibble(  
  x = 1:3,  
  y = list(1:2, 1:3, 1:4) # lists of different length on each subject  
)
```

```
## # A tibble: 3 x 2  
##       x y  
##   <int> <list>  
## 1     1 1 <int [2]>  
## 2     2 2 <int [3]>  
## 3     3 3 <int [4]>
```

## Logical and relational operators

# Logical operators

- ▶ The basic logical operators are described in `help("Logic")`.
- ▶ `!` is NOT
- ▶ `&` and `&&` are AND, with `&` acting vector-wise and `&&` acting on scalars
- ▶ `|` and `||` are OR, with `|` acting vector-wise and `||` acting on scalars
- ▶ Make sure you understand the following:

```
x <- c(TRUE,TRUE,FALSE); y <- c(FALSE,TRUE,TRUE)
!x ; x&y ; x&& y ; x|y ; x||y
```

```
## [1] FALSE FALSE TRUE
```

```
## [1] FALSE TRUE FALSE
```

```
## [1] FALSE
```

```
## [1] TRUE TRUE TRUE
```

```
## [1] TRUE
```

- ▶ Notice how `&&` and `||` act on the first element of the vectors `x` and `y` and ignore all the rest.



# Relational operators

- ▶ Relational operators can be used to compare values in atomic vectors
  - ▶ See `help("Comparison")`
- ▶ `>` is greater than, `>=` is greater than or equal
- ▶ `<` is less than, `<=` is less than or equal
- ▶ `==` is equal and `!=` is not equal
- ▶ Make sure you understand the following:

```
x <- 1:3; y <- 3:1  
x>y ; x>=y ; x<y ; x<=y ; x==y ; x!=y
```

```
## [1] FALSE FALSE  TRUE
```

```
## [1] FALSE  TRUE  TRUE
```

```
## [1]  TRUE FALSE FALSE
```

```
## [1]  TRUE  TRUE FALSE
```

```
## [1] FALSE  TRUE FALSE
```

```
## [1]  TRUE FALSE  TRUE
```

# Subsetting vectors with logical expressions

- Can subset with logicals and [:

```
avec
```

```
##    age weight height  
##    54    210     77
```

```
avec>100
```

```
##    age weight height  
## FALSE    TRUE  FALSE
```

```
avec[avec>100]
```

```
## weight  
##    210
```

```
avec[avec>54 & avec<100]
```

```
## height  
##     77
```

## Subsetting matrices with logical expressions

- ▶ Can also subset matrices, but results may not be as expected:

```
A
```

```
##          var1 var2
## subj1      1    2
## subj2      3    4
```

```
A>1
```

```
##          var1 var2
## subj1 FALSE TRUE
## subj2  TRUE TRUE
```

```
A[A>1] # coerces to a vector
```

```
## [1] 3 2 4
```

# Subset and assign with logical expressions

- ▶ Combine subset and assign to change the value of objects

```
A[A>1] <- 9
```

```
A
```

```
##      var1 var2
```

```
## subj1    1    9
```

```
## subj2    9    9
```

- ▶ In the above substitution, the vector 9 is shorter than the three elements in A>1 so R “recycles” the 9 three times.

## Be careful about recycling:

```
A[A>1] <- c(-10,10) # Throws a warning
```

```
## Warning in A[A > 1] <- c(-10, 10): number of items to replace is not  
## of replacement length
```

```
A # R used c(-10,10), then just the -10
```

```
##      var1 var2  
## subj1    1   10  
## subj2  -10  -10
```

## Aside: Special values

# Missing values

- ▶ R has a special data code for missing data: NA
- ▶ Test for and set missing values with `is.na()`

```
avec
```

```
##      age weight height  
##      54     210      77
```

```
is.na(avec)
```

```
##      age weight height  
## FALSE  FALSE  FALSE
```

```
is.na(avec) <- 2  
avec
```

```
##      age weight height  
##      54      NA      77
```

# Infinite and undefined values

- ▶ R has a special codes for infinite values (Inf) and undefined values (NaN).
- ▶ Test for Inf and NaN with `is.infinite()` and `is.nan()`.

```
ii <- 1/0 ; nn <- 0/0
```

```
ii
```

```
## [1] Inf
```

```
is.infinite(ii)
```

```
## [1] TRUE
```

```
nn
```

```
## [1] NaN
```

```
is.nan(nn)
```

```
## [1] TRUE
```



# The null object

- ▶ The null object, NULL, is an un-typed no-value object.
  - ▶ Test for NULL with `is.null()`
  - ▶ NULL can be used to initialize objects that will be created through combining, rbinding, etc.

```
x <- NULL; is.null(x)
```

```
## [1] TRUE
```

```
x <- c(x,1); x <- c(x,2); x
```

```
## [1] 1 2
```

```
# etc., or as a loop (more on these later)
```

```
x <- NULL  
for(i in 1:2) {  
  x <- c(x,i)  
}  
x
```

```
## [1] 1 2
```